

# General concepts in supervised learning

---

- ❑ Training vs. testing
- ❑ Loss function for training
  - Regression: RSS or MSE for regression
  - Classification: Log Likelihood or Log posterior probability, cross entropy
  - Regularization: constrain the model parameters based on some prior knowledge
- ❑ Performance metric (for test samples)
  - Regression: RSS
  - Classification: Accuracy, confusion matrix, sensitivity, specificity, precision, recall, ROC curve, AUC
- ❑ Cross validation to estimate the test performance
- ❑ Cross validation for model selection or hyper parameter optimization
- ❑ K>2 folds are needed when we have limited data
- ❑ Bagging (model averaging)

# Regression methods

---

- Linear regression (linear in model parameters):  $\hat{y}_i = \sum_{j=0}^p A_{ij}\beta_j$ 
  - Least squares fitting:  $\text{RSS} = \sum_{i=1}^n (y_i - \hat{y}_i)^2 \rightarrow \min: \boldsymbol{\beta} = (A^T A)^{-1} A^T \mathbf{y}$
  - Should normalize the data
- Regularization:
  - Ridge:  $J(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j|^2$  (favor small coefficients)
  - LASSO:  $J(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \alpha \sum_{j=1}^d |\beta_j|$  (favor sparse set of coefficients, many are zero)
    - Can be used for feature selection
    - Determine  $\alpha$  through cross validation
- Support vector regression
  - Did not discuss in class, but very powerful especially with nonlinear kernel
- Neural net regression
- Decision tree and random forest for regression

# Classification methods

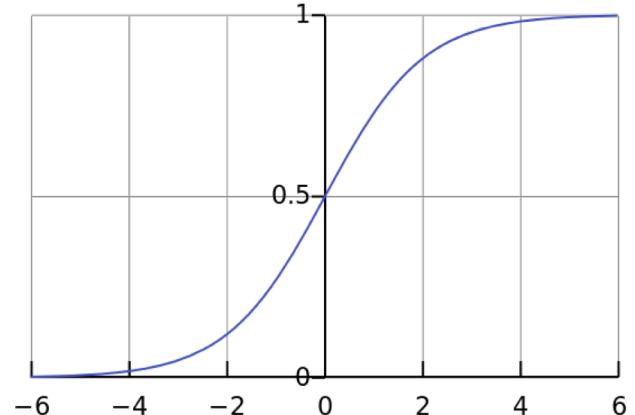
---

- Logistic regression
- Support vector classifier
- Neural net: fully connected
- Neural net: convolution layers + fully connected
- Decision trees and random forest

# Logistic regression

## ❑ Binary classification)

- Linear discriminant function  $z = w_0 + \sum_{j=1}^k w_k x_k$
- Mapping  $z$  to probability using sigmoidal:  $f(z) = 1/(1 + e^{-z})$
- Minimize log likelihood = binary cross entropy
- A linear classifier as it divides the feature space using hyperplane
- **Good only if the data are approximately linearly separable** ☹
- Need to transform original features if not linearly separable



## ❑ Multiclass

- $z_k = \mathbf{w}_k^T \mathbf{x} + w_{0k}$
- $g_k(\mathbf{z}) = \frac{e^{z_k}}{\sum_{\ell=1}^K e^{z_\ell}}$  (softmax)
- Train using multi-class cross entropy

# Support vector machine

- ❑ Also use linear discriminant (if using linear kernel):

$$z_i = b + \mathbf{w}^T \mathbf{x}_i$$

- ❑ Minimize a weighted average of the hinge loss (No loss if within the margin) and the margin. C chosen by cross validation

$$C \sum_{i=1}^N \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b)) + \frac{1}{2} \|\mathbf{w}\|^2$$

- ❑ Support vectors: samples within the margin or on the wrong side of the line

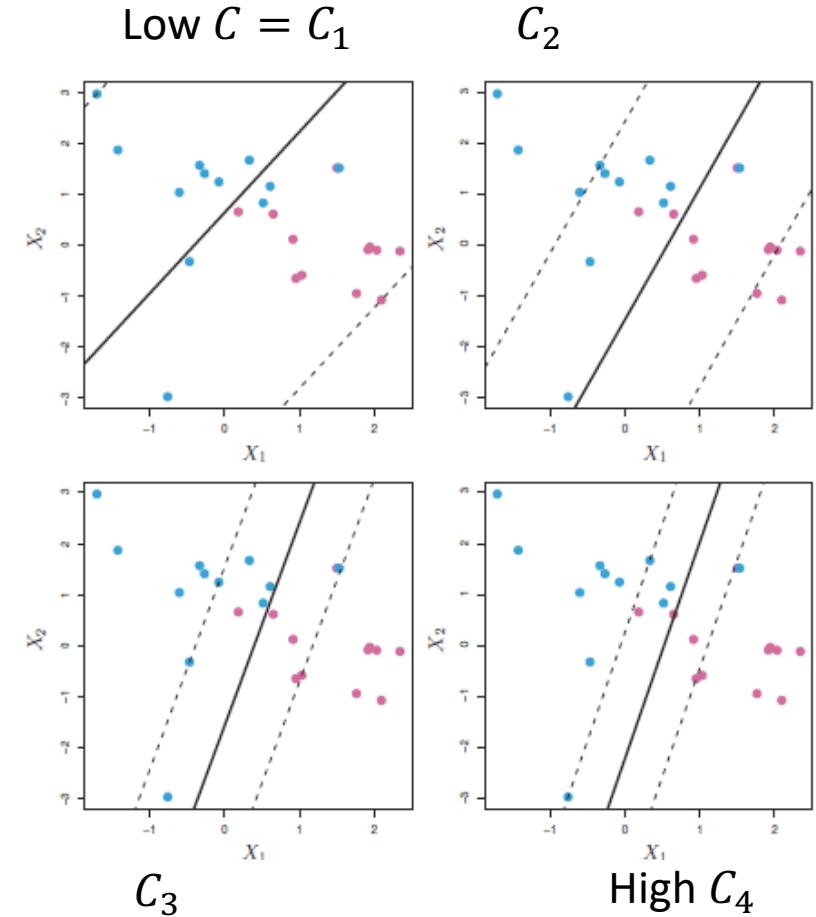
- $\mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i$  ( $\mathbf{x}_i$  are support vectors)

- $\hat{z}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$  (weighted average of  $y_i$  with weights proportional to "correlation"  $\mathbf{x}_i^T \mathbf{x}$ .)

- ❑ Can be extended to nonlinear partition by using non-linear kernels

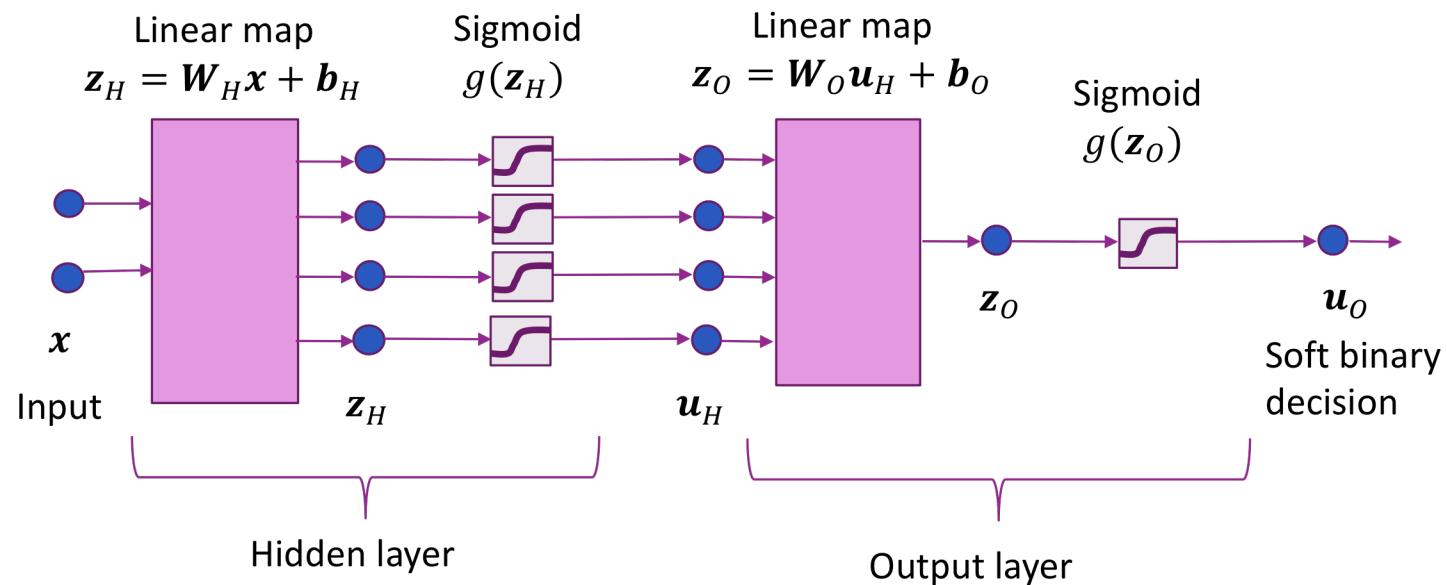
$$z = b + \mathbf{w}^T \phi(\mathbf{x}) = b + \sum_{i=1}^N \alpha_i y_i \phi(\mathbf{x}_i)^T \phi(\mathbf{x})$$

- ❑ Need to save many support vectors 😞



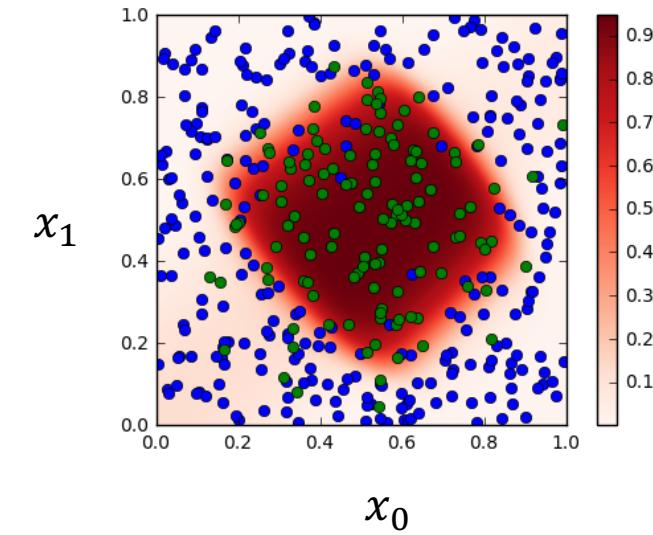
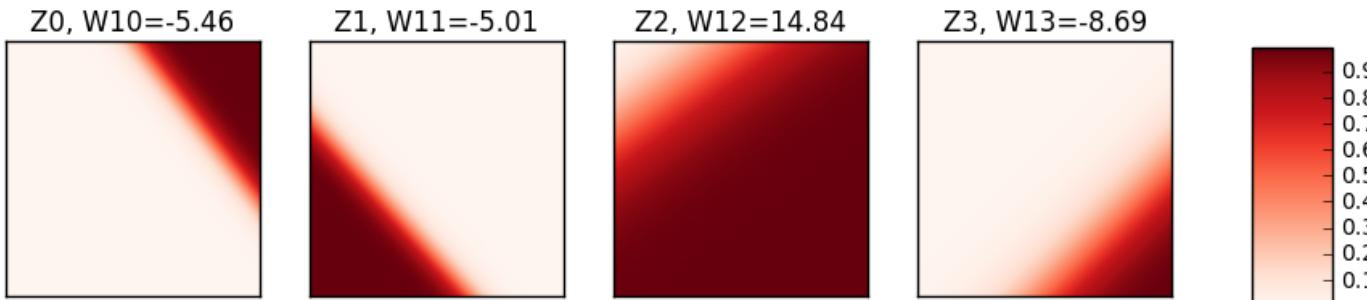
# Neural networks

- Stacks of logistic regression layers
- Nonlinear mapping after each linear combination is important!
- In principle, two layers with sufficient number of hidden nodes can realize any function



# Step 1 Outputs and Step 2 Outputs

---

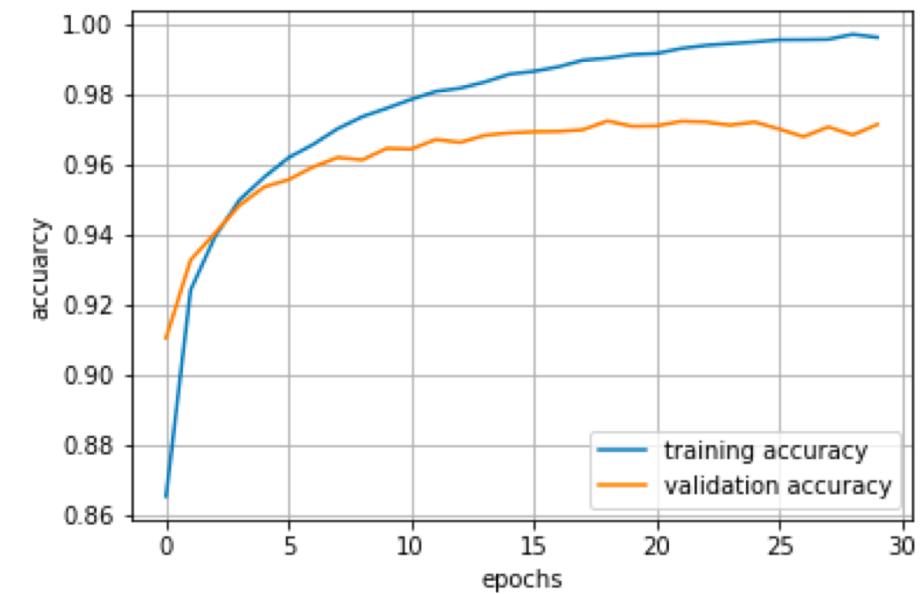


- ❑ Each output from step 1 is from a linear classifier with soft decision (Logistic regression)
- ❑ Final output is a weighted average of step 1 outputs using the weights indicated on top of the figures

# Training a neural net

---

- Minimize either RSS (for regression) or cross entropy (for classification) plus some regularization term
- Optimize parameters using gradient descent: Chain rule -> error backpropagation
- Stochastic gradient descent
  - Batches, epochs
  - Looking at the loss curves (training and validation) to determine when to stop
- Using Keras
  - Step 1. Describe model architecture
    - Number of hidden units, output units, activations, ...
  - Step 2. Select an optimizer
  - Step 3. Select a loss function and compile the model
  - Step 4. Fit the model
  - Step 5. Test / use the model
  - Need to know how to organize your data and labels in tensors



# Gradients on a Computation Graph

- ❑ **Backpropagation:** Compute gradients backwards

- Use tensor dot products and chain rule

- ❑ First compute all derivatives of all the variables

- $\partial L / \partial z_O$
  - $\partial L / \partial u_H = \langle \partial L / \partial z_O, \partial z_O / \partial u_H \rangle$
  - $\partial L / \partial z_H = \langle \partial L / \partial u_H, \partial u_H / \partial z_H \rangle$

- ❑ Then compute gradient of parameters:

- $\partial L / \partial W_O = \langle \partial L / \partial z_O, \partial z_O / \partial W_O \rangle$
  - $\partial L / \partial b_O = \langle \partial L / \partial z_O, \partial z_O / \partial b_O \rangle$
  - $\partial L / \partial W_H = \langle \partial L / \partial z_H, \partial z_H / \partial W_H \rangle$
  - $\partial L / \partial b_H = \langle \partial L / \partial z_H, \partial z_H / \partial b_H \rangle$

- ❑ You should know how to do this for a 2 layer network

